

CPSC 213: Introduction to Computer Systems

Unit 1a: Numbers and Memory

Jordon Johnson

Adapted from slides by Jonatan Schroeder, Mike Feeley, and Robert Xiao

Overview

- Reading
 - Companion: 2.2.2
 - Textbook: 2.1 - 2.3
- Learning Objectives
 - know the number of bits in a byte and the number of bytes in a short, long and quad
 - determine whether an address is aligned to a given size
 - translate between integers and values stored in memory for both big- and little-endian machines
 - evaluate and write Java expressions using bitwise operators (&, |, <<, >>, and >>>)
 - determine when sign extension is unwanted and eliminate it in Java
 - evaluate and write C expressions that include type casting and the addressing operators (& and *)
 - translate integer values by hand (no calculator) between binary and hexadecimal, add/subtract hexadecimal numbers and convert small numbers between binary and decimal

A Simple Computing Machine

- **Memory**


- stores data encoded as bits
- program instructions and state (variables, objects, etc.)

- **CPU**

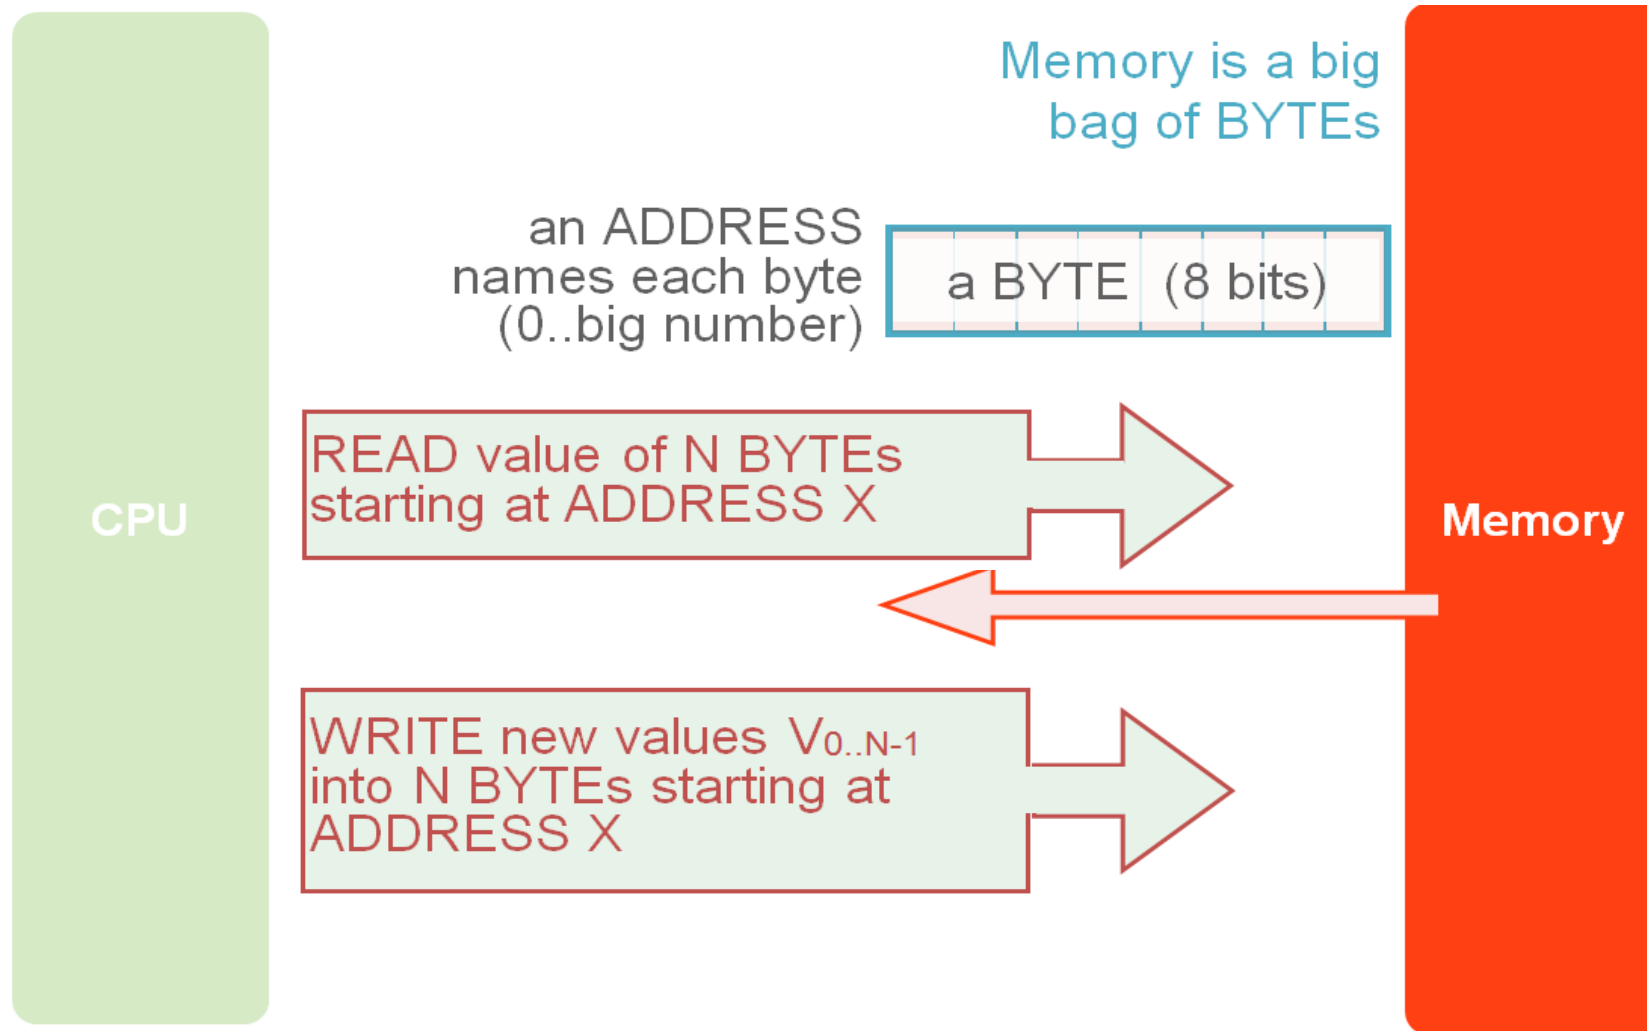
- reads instruction and data from memory
- performs specified computation and writes result back to memory

- **Example**

- $C = A + B$
- memory stores: add instruction, and variables A, B and C
- CPU reads instruction and values of A and B, adds values, writes result to C




Memory



Memory

This is an **implementation** of memory; we will focus on the **abstraction**



Memory

- **Naming**

- unit of addressing is a **byte** (8 bits)
- every byte of memory has a **unique address**
- some machines have 32-bit addresses, some have 64-bit addresses
 - We will (usually) assume our machine uses 32-bit addresses, and that all addresses are valid

- **Access**

- many things are too big to fit in a single byte
 - unsigned numbers > 255 , signed numbers < -128 or > 127 , most instructions, etc.
- CPU accesses memory in contiguous, power-of-two-size chunks of bytes
- address of a chunk is address of its **first byte**

Memory

Integer Data Types by Size

# bytes	# bits	C	Java	Asm
1	8	char	byte	b byte
2	16	short	short	w word
4	32	int	int	l long
8	64	long	long	q quad

We will use only 32-bit integers

Numbers and Representation

- Sometimes we are interested in the **integer value** of a chunk of bytes
 - base 10, **decimal**, is commonly used to represent this number (our “normal” number system)
 - we need to **convert** from binary to decimal to get this value
- Sometimes we are more interested in **bits** themselves
 - In such cases the decimal value isn’t particularly important
 - For example, consider **memory addresses**
 - big numbers that name power-of-two size things
 - we do not usually care what the base-10 value of an address is
 - we’d like a power-of-two sized way to name addresses

Numbers and Representation

- We might use base-2, **binary**
 - a small 256-byte memory has addresses 0_2 to 11111111_2
 - may be represented as **0b**11111111
 - becomes tedious and hard to read as addresses get larger
- Once we used base-8, **octal**
 - 64-KB memory addresses go up to $1111111111111111_2 = 177777_8$
 - may be represented as **0o**177777
 - gets tedious and hard to read too
- Now we use base-16, **hexadecimal**
 - 4-GB memory addresses go up to $377777777777_8 = \text{ffffffff}_{16}$
 - if you don't have subscripts, ffffffff_{16} is written as **0x**ffffffff

Binary ⇔ Hexadecimal

01101010010101010000111010100011

- 4 bits in a hex “digit”, a hexit (or “nibble”)

0110 1010 0101 0101 0000 1110 1010 0011

- Consider ONE hexit at a time

6 a 5 5 0 e a 3

0x6a550ea3

- A byte (8 bits) is just 2 hexits ($2^8=16^2$):

0x6a550ea3 => 0x6a 0x55 0x0e 0xa3

hex	bin
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111
8	1000
9	1001
A	1010
B	1011
C	1100
D	1101
E	1110
F	1111

Which of the following statements is true?

- A. The Java constants 16 and 0x10 are exactly the same integer
- B. The Java constants 16 and 0x10 are different integers
- C. Neither of the statements above is always true
- D. I don't know
- E. 42 is the answer to the ultimate question of life, the universe, and everything


Hexadecimal Operations

- We use hexadecimal for addresses
 - We don't really care what their base-10 value is
- Addition/subtraction in hex
 - You could convert both to decimal, but that might be too tedious
 - You can **calculate in hex directly**
 - Alternative for subtraction: use addition with **two's complement**
- Remember:
 - carry when result is $0x10 == 16_{10}$ or more
 - hexits A..F convert to their decimal value

hex	bin
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111
8	1000
9	1001
A	1010
B	1011
C	1100
D	1101
E	1110
F	1111


Object A is at address 0x10d4, and object B at 0x1110. They are stored contiguously in memory (i.e., they are adjacent to each other).

How big is A (in bytes)?



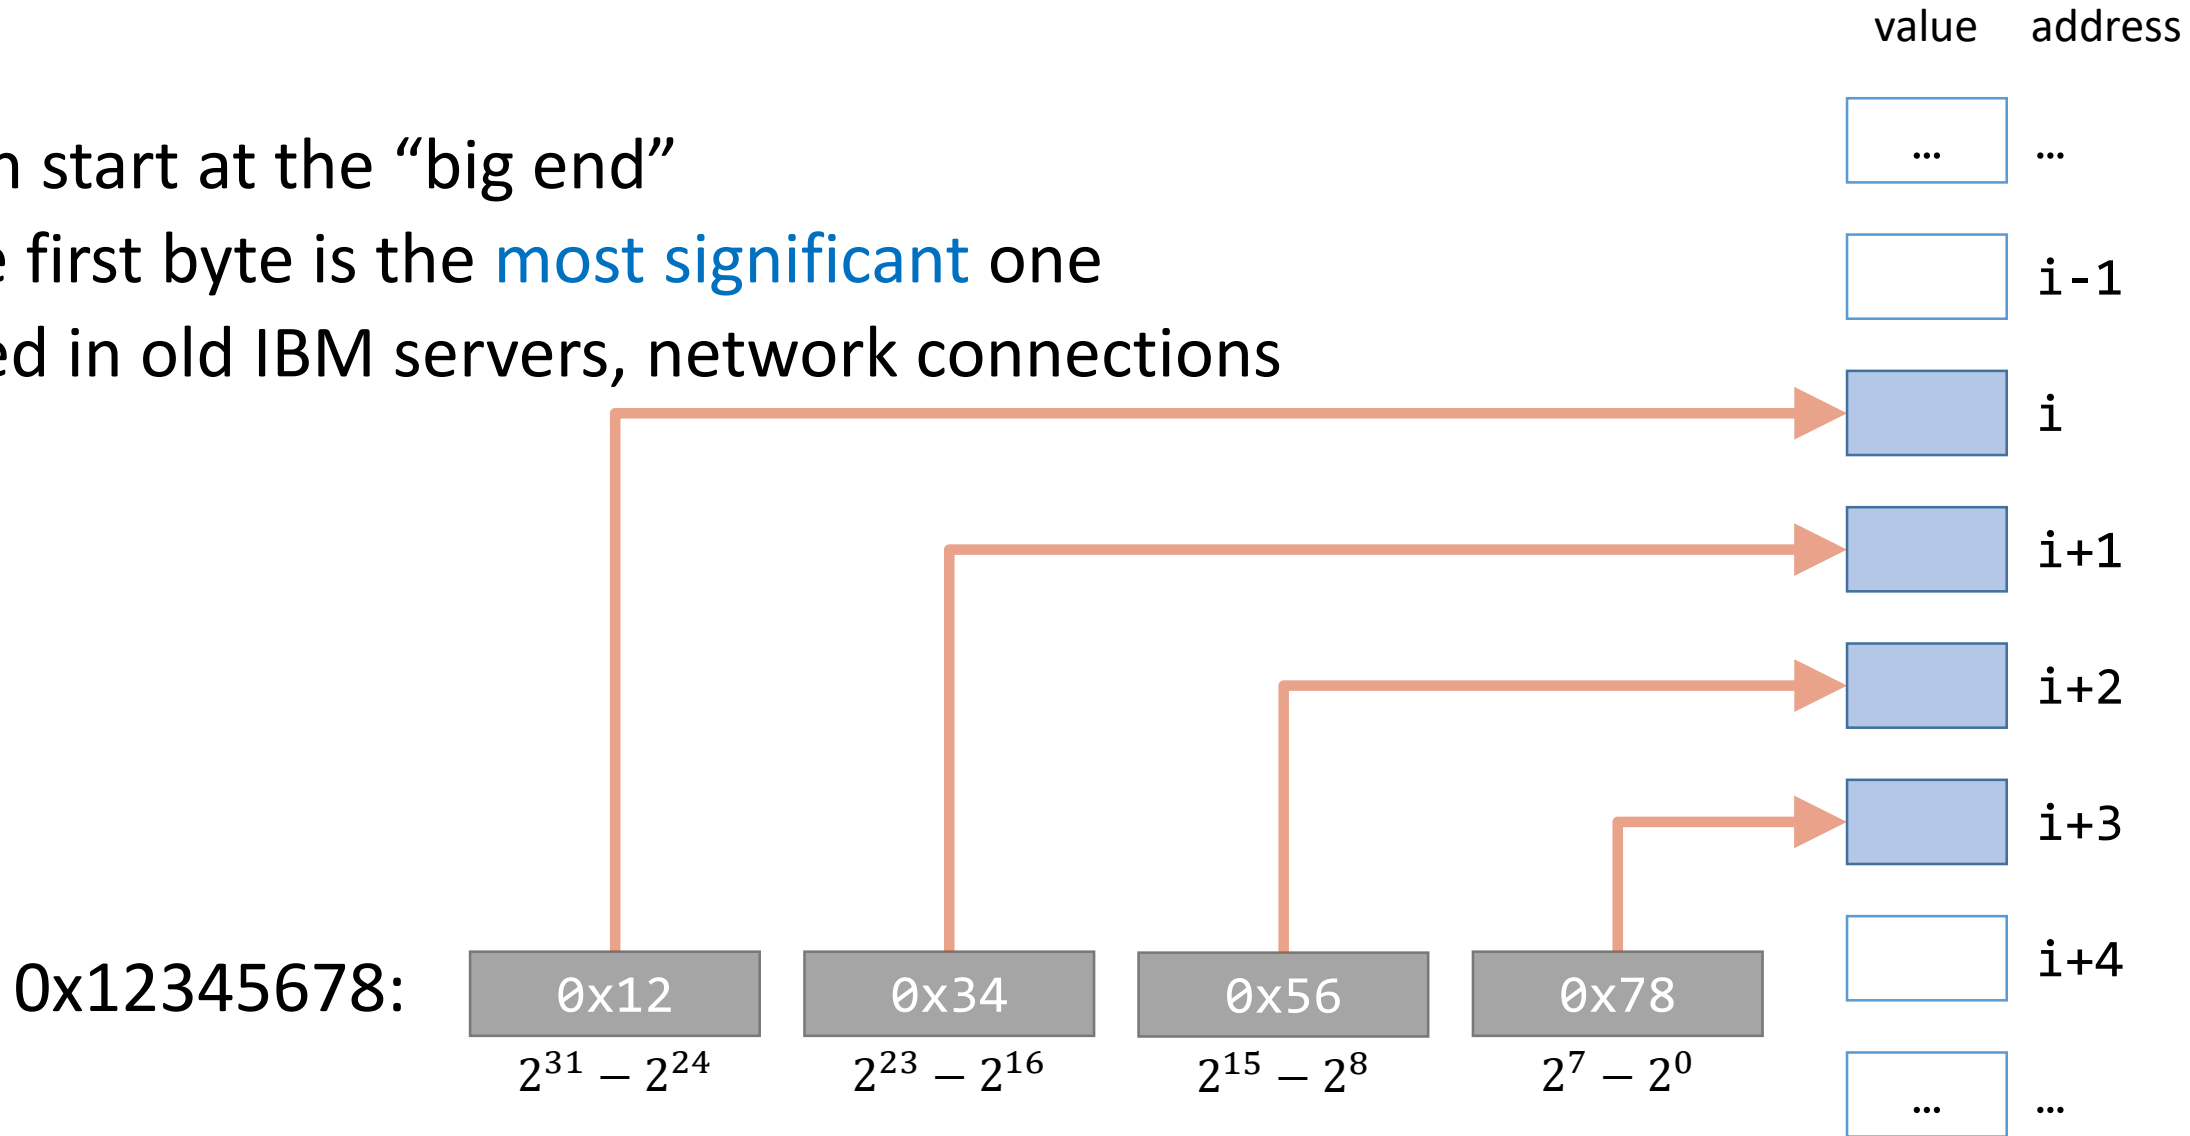
Making Integers from Bytes

- First architectural design decision:
 - assembling memory bytes into integers
- Consider a 32-bit integer (`int` in Java or C)
 - It uses 4 bytes
 - If memory address is `i`, then we also need bytes at `i+1`, `i+2`, `i+3`
 - Example: if address is `0x2014`, then integer is in `0x2014`, `0x2015`, `0x2016`, `0x2017`
- What do each of these bytes represent?



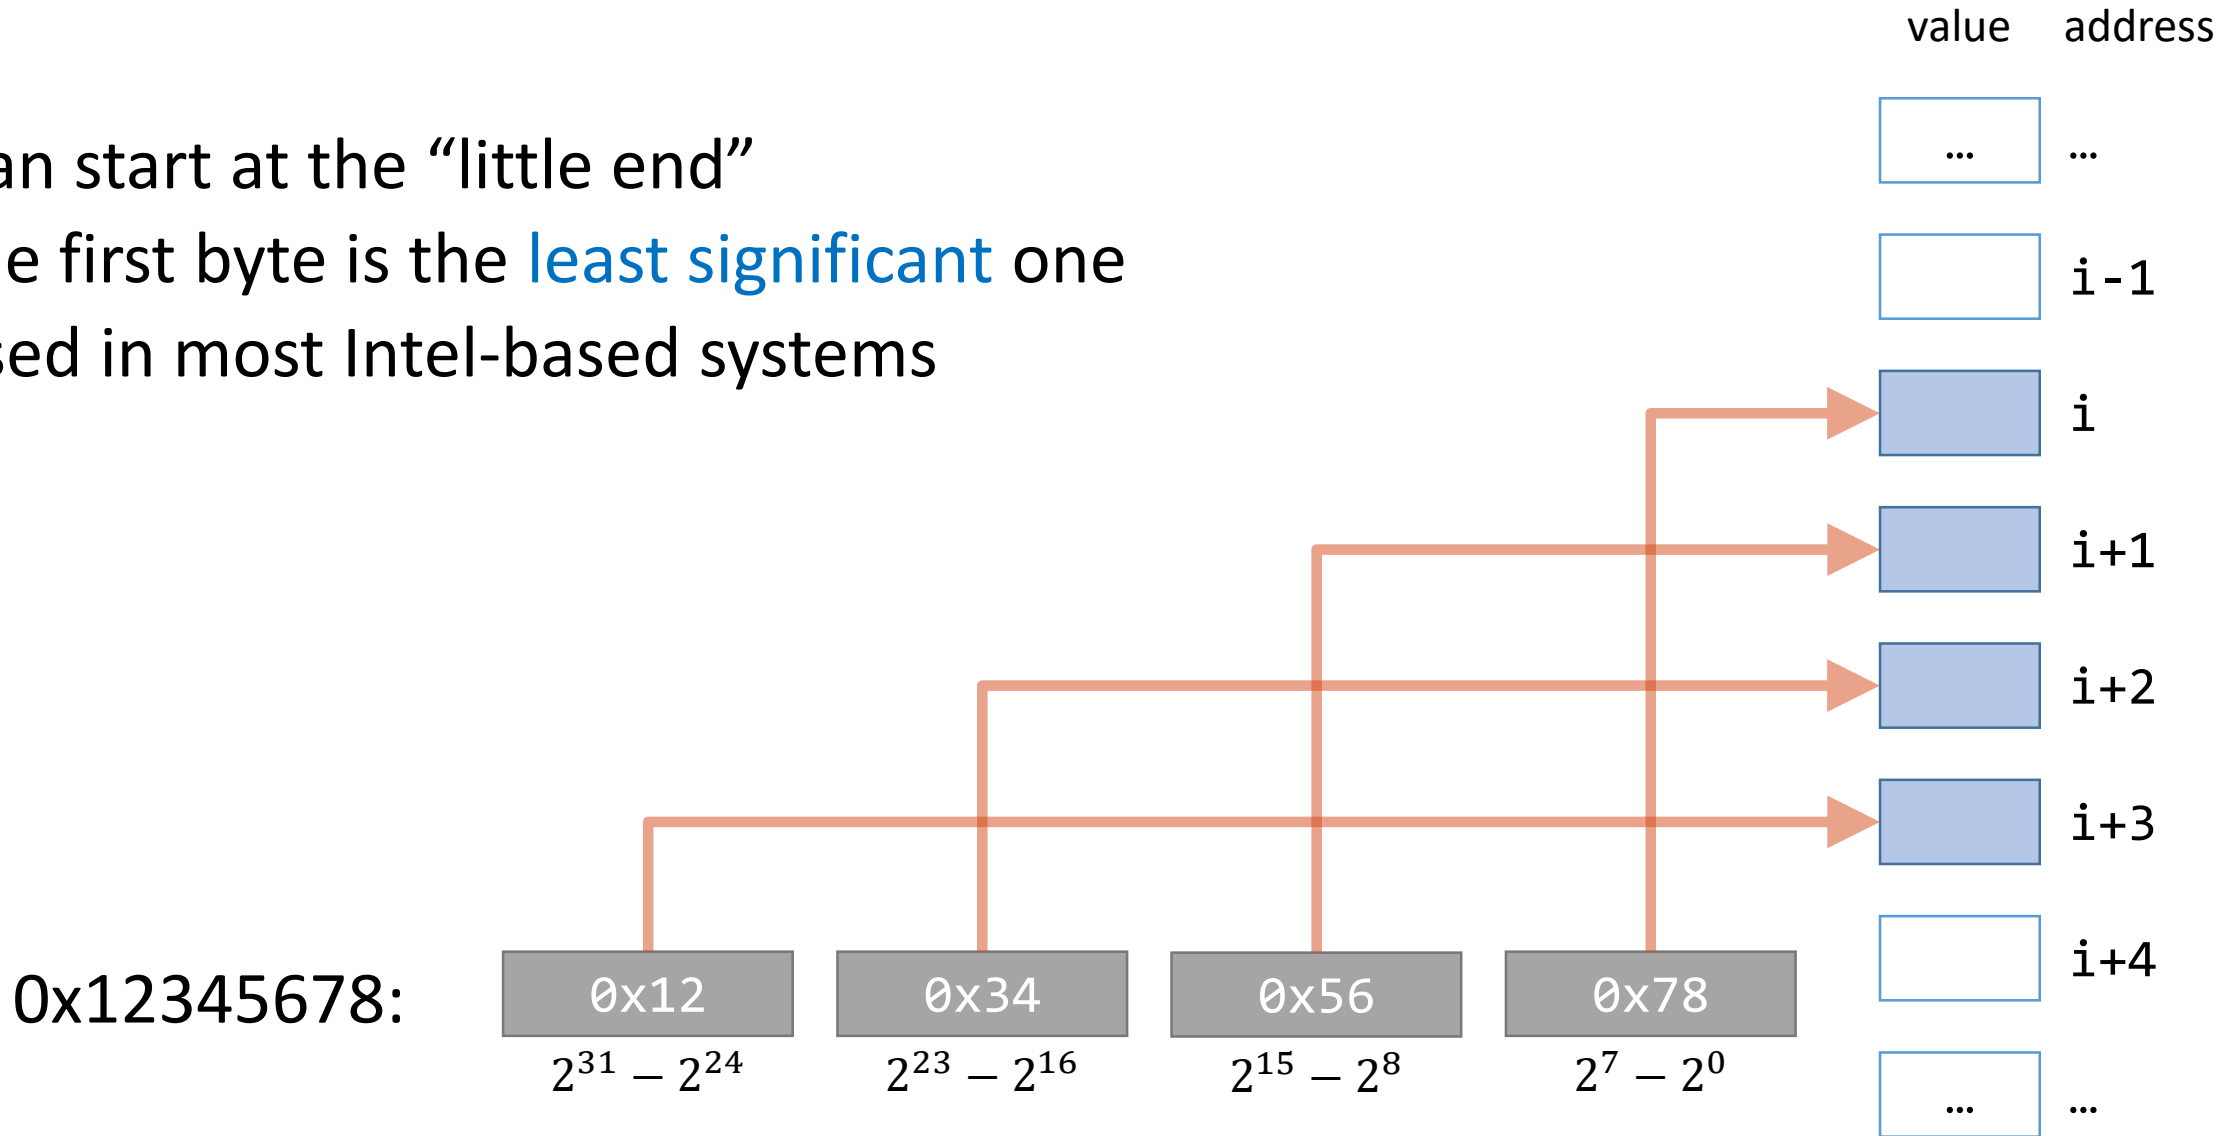
Big-Endian

- We can start at the “big end”
 - The first byte is the **most significant** one
 - Used in old IBM servers, network connections



Little-Endian

- We can start at the “little end”
 - The first byte is the **least significant** one
 - Used in most Intel-based systems



The memory of some machines stores Big Endian integers.

- A. True
- B. False


What is the Little-Endian 4-byte integer value at address 0x4?

- A. 0xc1406b37
- B. 0x1c04b673
- C. 0x73b6041c
- D. 0x376b40c1
- E. 0x739a8732


Address	Value
0x0:	0xfe
0x1:	0x32
0x2:	0x87
0x3:	0x9a
0x4:	0x73
0x5:	0xb6
0x6:	0x04
0x7:	0x1c

Should We Put Data Just Anywhere?

- We could place a 4-byte integer at any address



- However, requiring addresses to be **aligned** is better for hardware



Address Alignment

- What is an “aligned” address?
 - Address whose numeric value is a **multiple** of the object **size**
 - It depends on the object; it gets slightly more complicated with **arrays** and **structs (spoilers!)**
- Aligned addresses are better:
 - You can fit **two shorts** inside an **int**, etc.
 - This is significant in **arrays** as well
 - Some CPUs don't support misaligned addresses
 - Intel: misaligned access is supported but slower

Address Alignment

- CPU implementation encourages alignment
 - Memory is organized internally into chunks (**blocks**)
 - Every **memory access** requires accessing a **whole block**
 - *Details: see CPSC 313*
- CPU memory access looks like:
 - Read/write N bytes starting at address A
- This is translated by memory to:
 - Block B contains addresses $X..(X + \text{blocksize} - 1)$
 - $X \leq A \leq X + \text{blocksize} - 1$
 - So A is at some offset (let's call it O) from the beginning of block B
 - Read/write N bytes starting at O^{th} byte of block B ($A = X + O$)

 O for offset, not zero

Address Alignment

- (From last slide):
 - Block B contains addresses $X..(X + \textit{blocksize} - 1)$
 - $X \leq A \leq X + \textit{blocksize} - 1$
 - Read/write N bytes starting at O^{th} byte of block B ($A = X + O$)
- How is this **simplified** if:
 - $\textit{blocksize}$ is a power of 2,
 - N is a power of 2,
 - A is aligned to N ?

Which of the following statements is/are true?

- A. the address $6_{10} = 110_2$ is aligned for a short (2 byte) integer
- B. the address $6_{10} = 110_2$ is aligned for a 4 byte integer
- C. the address $20_{10} = 10100_2$ is aligned for a long (8 byte) integer
- D. Two or more of the above
- E. None of the above

Which of the following statements is (are) true?

- A. The address 0x14 is aligned for addressing a 2-byte integer, but not a 4-byte integer
- B. The address 0x14 is aligned for addressing a 2-byte or a 4-byte integer, but not an 8-byte integer
- C. The address 0x14 is aligned for addressing a 2-byte, 4-byte, or 8-byte integer
- D. None of the above
- E. If I say the answer is 42, maybe I'll get the point anyway

hex	bin
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111
8	1000
9	1001
A	1010
B	1011
C	1100
D	1101
E	1110
F	1111

Changing Data Types: Extending

- [illegible]

Changing Data Types: Truncating

- Truncating an integer: reducing the number of bytes used to store it
`int i = -6; // stored as 1111111111111111111111111111111111111010`
`byte b = i; // stored as 111010`
- What could go wrong?
 - What value would b get if i were 256? or 128?
- Java warns you if you truncate implicitly
 - To avoid warning, cast explicitly:
`byte b = (byte) i;`

Bit Operations in C/Java

- $a \ll b$: shift all bits in a to the left b times, fill remaining right bits with zero
- $a \gg b$: shift all bits in a to the right b times
 - C: if a is unsigned, zero-extends, otherwise sign-extends
 - Java: operator \gg sign-extends, operator \ggg zero-extends
- $a \& b$: AND applied to corresponding bits in a and b
- $a | b$: OR applied to corresponding bits in a and b
- $a \wedge b$: XOR applied to corresponding bits in a and b
- $\sim a$: inverts every bit of a

Making Use of Bit Operations

- Shifting multiplies/divides by power of 2
 - “ $a \ll b$ ” is equivalent to $a \times 2^b$
 - “ $a \gg b$ ” is equivalent to $a / 2^b$
- Example: 22 in binary is 00010110
 - 11 is 00001011 (22 shifted right once, $22 / 2^1$)
 - 44 is 00101100 (22 shifted left once, 22×2^1)
 - 88 is 01011000 (22 shifted left twice, 22×2^2)
- Works for negative numbers too, if using sign-extended shift
 - -22 is 11101010
 - -11 is 11110101 (-22 shifted right once, $-22 / 2^1$)
 - -44 is 11010100 (-22 shifted left once, -22×2^1)
 - -88 is 10101000 (-22 shifted left twice, -22×2^2)

Making Use of Bit Operations

- Let's use our example from before:

```
byte b = -6; // stored as 11111010
```

```
int i = b; // stored as 1111...1111 11111010
```

- How can we get it to **zero-extend** instead of **sign-extend**?

- **Answer:** using **bit operations**:

```
// 0xFF in bits is: 0000...000011111111
```

```
int i = b & 0xFF; // stored as 0000...000011111010
```

- Note that **the result is different**: 250 instead of -6
 - That's because zero-extension is used for unsigned integers

What is the value of `i` after this Java statement executes?

```
int i = 0xff8b0000 & 0x00ff0000;
```

- A. `0xfffff0000`
- B. `0x0000008b`
- C. `0x008b0000`
- D. `0xff8b0000`
- E. None of the above

What is the value of `i` after this Java statement executes?

```
int i = 0x0000008b << 16;
```

- A. `0x008b`
- B. `0x0000008b`
- C. `0x008b0000`
- D. `0xff8b0000`
- E. None of the above

What is the value of `i` after this Java statement executes?

```
int i = 0x8b << 16;
```

- A. `0x8b`
- B. `0x0000008b`
- C. `0x008b0000`
- D. `0xff8b0000`
- E. None of the above

What is the value of `i` after this Java statement executes?

```
int i = ((byte) 0x8b) << 16;
```

- A. `0x8b`
- B. `0x0000008b`
- C. `0x008b0000`
- D. `0xff8b0000`
- E. None of the above